

SAE fidelity: read the numbers, then the color
Before-word geometry excludes BOS prefixes; behavioral checks use four sampled windows per corpus by default

How much of each state changes?



How much do predictions change?

